

ds-exercise1-b1-ui22cs03

August 9, 2024

1 Excercise 1 by UI22CS03

1.1 Objectives

- Import Libraries
 - Lab Exercises
 - Identifying duplicates
 - Plotting Scatterplots
-

1.2 Import Libraries

Import the libraries we need

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Read the csv file

```
[3]: data = pd.read_csv("TeachingRatings_ui22cs03.csv")
print(data)
```

	rownames	minority	age	gender	credits	beauty	eval	division	native	\
0	1	yes	36	female	more	0.289916	4.3	upper	yes	
1	2	no	59	male	more	-0.737732	4.5	upper	yes	
2	3	no	51	male	more	-0.571984	3.7	upper	yes	
3	4	no	40	female	more	-0.677963	4.3	upper	yes	
4	5	no	31	female	more	1.509794	4.4	upper	yes	
..	
458	459	no	32	male	more	1.231394	3.2	lower	yes	
459	460	no	32	male	more	1.231394	4.3	upper	yes	
460	461	yes	42	female	more	0.420400	3.3	upper	no	
461	462	yes	42	female	more	0.420400	3.2	upper	no	
462	463	yes	42	female	single	0.420400	4.1	lower	no	
	tenure	students	allstudents	prof						
0	yes	24	43	1						

1	yes	17	20	2
2	yes	55	55	3
3	yes	40	46	4
4	yes	42	48	5
..
458	yes	9	21	93
459	yes	52	86	93
460	yes	52	67	94
461	yes	54	66	94
462	yes	28	35	94

[463 rows x 13 columns]

1.3 Display information about the dataset

1. Structure of the dataframe
2. Describe the dataset
3. Number of rows and columns

```
[4]: # 1. Structure of the dataframe
print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 463 entries, 0 to 462
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   rownames         463 non-null    int64
1   minority         463 non-null    object
2   age              463 non-null    int64
3   gender           463 non-null    object
4   credits          463 non-null    object
5   beauty           463 non-null    float64
6   eval             463 non-null    float64
7   division         463 non-null    object
8   native           463 non-null    object
9   tenure           463 non-null    object
10  students         463 non-null    int64
11  allstudents      463 non-null    int64
12  prof             463 non-null    int64
dtypes: float64(2), int64(5), object(6)
memory usage: 47.1+ KB
None
```

	rownames	age	beauty	eval	students \
count	463.000000	463.000000	4.630000e+02	463.000000	463.000000
mean	232.000000	48.365011	6.263499e-08	3.998272	36.624190
std	133.800847	9.802742	7.886477e-01	0.554866	45.018481
min	1.000000	29.000000	-1.450494e+00	2.100000	5.000000

25%	116.500000	42.000000	-6.562689e-01	3.600000	15.000000
50%	232.000000	48.000000	-6.801430e-02	4.000000	23.000000
75%	347.500000	57.000000	5.456024e-01	4.400000	40.000000
max	463.000000	73.000000	1.970023e+00	5.000000	380.000000

	allstudents	prof
count	463.000000	463.000000
mean	55.177106	45.434125
std	75.072800	27.508902
min	8.000000	1.000000
25%	19.000000	20.000000
50%	29.000000	44.000000
75%	60.000000	70.500000
max	581.000000	94.000000

Number of rows and columns: (463, 13)

```
[5]: # 2. Describe the dataset
print(data.describe())
```

	rownames	age	beauty	eval	students \
count	463.000000	463.000000	4.630000e+02	463.000000	463.000000
mean	232.000000	48.365011	6.263499e-08	3.998272	36.624190
std	133.800847	9.802742	7.886477e-01	0.554866	45.018481
min	1.000000	29.000000	-1.450494e+00	2.100000	5.000000
25%	116.500000	42.000000	-6.562689e-01	3.600000	15.000000
50%	232.000000	48.000000	-6.801430e-02	4.000000	23.000000
75%	347.500000	57.000000	5.456024e-01	4.400000	40.000000
max	463.000000	73.000000	1.970023e+00	5.000000	380.000000

	allstudents	prof
count	463.000000	463.000000
mean	55.177106	45.434125
std	75.072800	27.508902
min	8.000000	1.000000
25%	19.000000	20.000000
50%	29.000000	44.000000
75%	60.000000	70.500000
max	581.000000	94.000000

```
[6]: # 3. Number of rows and columns
print("Number of rows and columns:", data.shape)
```

Number of rows and columns: (463, 13)

```
[7]: #Head Preview of data : Head means the first 5 datas
data.head()
```

```
[7]:      rownames minority  age  gender credits    beauty  eval division native \
0         1      yes   36  female    more  0.289916   4.3    upper    yes
1         2      no    59   male    more -0.737732   4.5    upper    yes
2         3      no    51   male    more -0.571984   3.7    upper    yes
3         4      no    40  female    more -0.677963   4.3    upper    yes
4         5      no    31  female    more  1.509794   4.4    upper    yes

      tenure  students  allstudents  prof
0    yes      24         43         1
1    yes      17         20         2
2    yes      55         55         3
3    yes      40         46         4
4    yes      42         48         5
```

```
[8]: #Tail Preview of data : Tail means the last 5 datas
data.tail()
```

```
[8]:      rownames minority  age  gender credits    beauty  eval division native \
458      459      no    32   male    more  1.231394   3.2    lower    yes
459      460      no    32   male    more  1.231394   4.3    upper    yes
460      461     yes    42  female    more  0.420400   3.3    upper    no
461      462     yes    42  female    more  0.420400   3.2    upper    no
462      463     yes    42  female  single  0.420400   4.1    lower    no

      tenure  students  allstudents  prof
458    yes      9         21         93
459    yes     52         86         93
460    yes     52         67         94
461    yes     54         66         94
462    yes     28         35         94
```

1.3.1 Identify all duplicate cases using prof variable - find the unique values of the prof variables

```
[9]: unique_prof = data['prof'].unique()
print(unique_prof)
```

```
[ 1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20 21 22 23 24
25 26 27 28 29 30 31 32 33 34 35 36 37 38 39 40 41 42 43 44 45 46 47 48
49 50 51 52 53 54 55 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71 72
73 74 75 76 77 78 79 80 81 82 83 84 85 86 87 88 89 90 91 92 93 94]
```

1.3.2 Print out the number of unique values in the prof variable

```
[11]: print(len(unique_prof))
```

1.3.3 Using all observations, Find the average and standard deviation for age

```
[12]: # Using all observations, Find the average and standard deviation for age
print("Average age:", data['age'].mean())
print("Standard deviation of age:", data['age'].std())
```

Average age: 48.365010799136066

Standard deviation of age: 9.802742037864817

1.3.4 Repeat the analysis by first filtering the data set to include one observation for each instructor with a total number of observations restricted to 94.

```
[21]: # Filter the dataset to include one observation per instructor
filtered_data = data.drop_duplicates(subset='prof')

# Restrict the number of observations to 94
filtered_data = filtered_data.head(94)

filtered_data.head()
```

```
[21]:
```

	rownames	minority	age	gender	credits	beauty	eval	division	native	\
0	1	yes	36	female	more	0.289916	4.3	upper	yes	
1	2	no	59	male	more	-0.737732	4.5	upper	yes	
2	3	no	51	male	more	-0.571984	3.7	upper	yes	
3	4	no	40	female	more	-0.677963	4.3	upper	yes	
4	5	no	31	female	more	1.509794	4.4	upper	yes	

	tenure	students	allstudents	prof
0	yes	24	43	1
1	yes	17	20	2
2	yes	55	55	3
3	yes	40	46	4
4	yes	42	48	5

```
[20]: filtered_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 94 entries, 0 to 93
Data columns (total 13 columns):
#   Column          Non-Null Count  Dtype
---  -
0   rownames        94 non-null    int64
1   minority        94 non-null    object
2   age             94 non-null    int64
3   gender          94 non-null    object
4   credits         94 non-null    object
5   beauty          94 non-null    float64
6   eval            94 non-null    float64
```

```

7   division      94 non-null    object
8   native         94 non-null    object
9   tenure         94 non-null    object
10  students       94 non-null    int64
11  allstudents    94 non-null    int64
12  prof           94 non-null    int64
dtypes: float64(2), int64(5), object(6)
memory usage: 10.3+ KB

```

Use the new dataset to get the mean and SD of age

```

[14]: # Calculate the average and standard deviation of age for the filtered dataset
print("Average age (filtered):", filtered_data['age'].mean())
print("Standard deviation of age (filtered):", filtered_data['age'].std())

```

```

Average age (filtered): 47.5531914893617
Standard deviation of age (filtered): 10.25651329515495

```

1.3.5 Using a bar chart, demonstrate if instructors teaching lower-division courses receive higher average teaching evaluations.

Find the average teaching evaluation in both groups of upper and lower-division

Plot the barplot

```

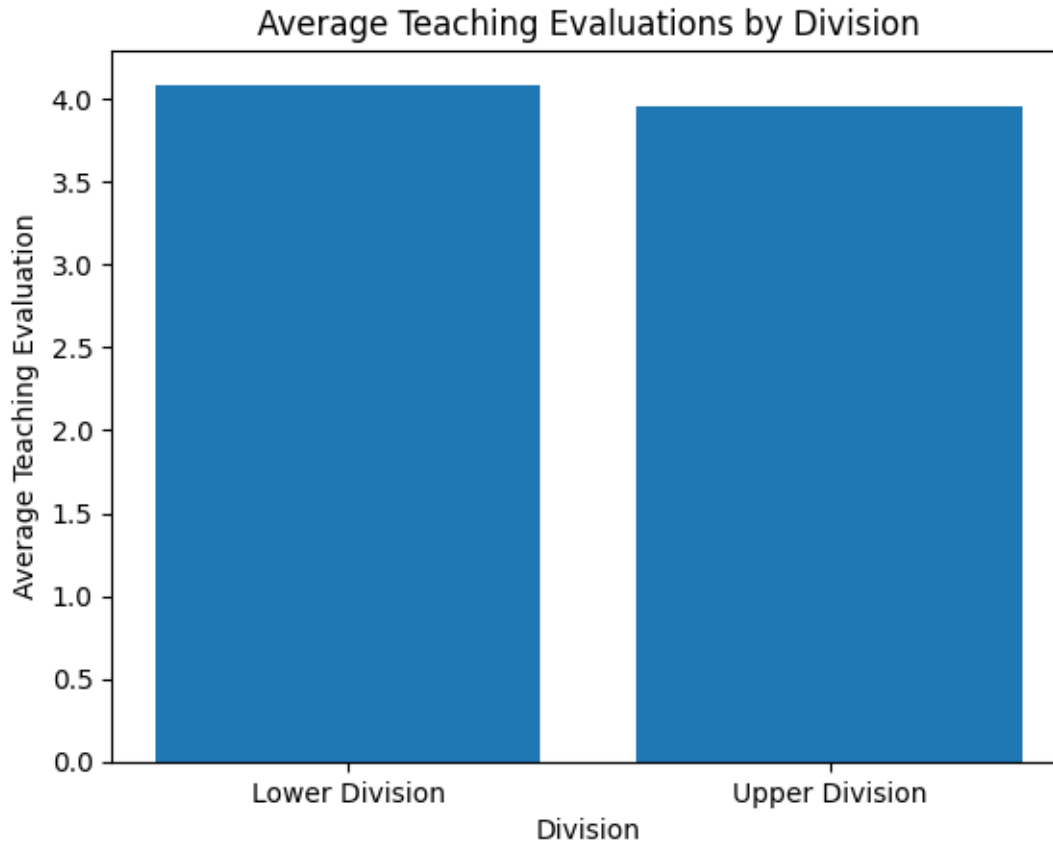
[22]: # Find the average teaching evaluation in both groups of upper and
      ↪ lower-division

# Calculate the average teaching evaluation for lower-division courses
lower_division_mean = data[data['division'] == 'lower']['eval'].mean()

# Calculate the average teaching evaluation for upper-division courses
upper_division_mean = data[data['division'] == 'upper']['eval'].mean()

# Create a bar chart
plt.bar(['Lower Division', 'Upper Division'], [lower_division_mean,
      ↪ upper_division_mean])
plt.xlabel('Division')
plt.ylabel('Average Teaching Evaluation')
plt.title('Average Teaching Evaluations by Division')
plt.show()

```



1.3.6 Using gender-differentiated scatter plots, plot the relationship between age and teaching evaluation scores.


```
[23]: # Create separate DataFrames for male and female instructors
male_instructors = data[data['gender'] == 'male']
female_instructors = data[data['gender'] == 'female']

# Create scatter plots for male and female instructors
plt.figure(figsize=(10, 6))

plt.scatter(male_instructors['age'], male_instructors['eval'], label='Male',
            alpha=0.7)
plt.scatter(female_instructors['age'], female_instructors['eval'],
            label='Female', alpha=0.7)

plt.xlabel('Age')
plt.ylabel('Teaching Evaluation Score')
plt.title('Relationship between Age and Teaching Evaluation Scores by Gender')
plt.legend()
```

```
plt.grid(True)
plt.show()
```



1.3.7 What is the number of courses taught by gender?

```
[26]: # Group the data by gender and count the number of courses
courses_by_gender = data.groupby('gender')['division'].count()

# Print the results
print(courses_by_gender)
```

```
gender
female    195
male      268
Name: division, dtype: int64
```

1.3.8 Create a group histogram of taught by gender and tenure

```
[27]: # Group the data by gender and tenure, then count the number of courses
courses_by_gender_tenure = data.groupby(['gender', 'tenure'])['division'].
    .count().unstack()

# Create a grouped histogram
courses_by_gender_tenure.plot(kind='bar', figsize=(10, 6))
```



```
plt.xlabel('Gender')
plt.ylabel('Number of Courses')
plt.title('Number of Courses Taught by Gender and Tenure')
plt.legend(title='Tenure')
plt.show()
```



1.3.9 Does age affect teaching evaluation scores?

```
[28]: # Calculate the correlation coefficient between age and teaching evaluation
      ↪ scores
correlation = data['age'].corr(data['eval'])

print("Correlation between age and teaching evaluation scores:", correlation)

# Create a scatter plot to visualize the relationship
plt.figure(figsize=(10, 6))
sns.scatterplot(x='age', y='eval', data=data)
plt.xlabel('Age')
plt.ylabel('Teaching Evaluation Score')
plt.title('Relationship between Age and Teaching Evaluation Scores')
plt.grid(True)
plt.show()
```

```


# Perform a simple linear regression to model the relationship
import statsmodels.api as sm

X = data['age']
y = data['eval']
X = sm.add_constant(X) # Add a constant to the model

model = sm.OLS(y, X).fit()
print(model.summary())

```

Correlation between age and teaching evaluation scores: -0.051696190877182864



OLS Regression Results

```

=====
Dep. Variable:          eval    R-squared:                0.003
Model:                  OLS     Adj. R-squared:           0.001
Method:                 Least Squares   F-statistic:             1.235
Date:                   Fri, 09 Aug 2024   Prob (F-statistic):      0.267
Time:                   07:21:20   Log-Likelihood:          -383.13
No. Observations:      463     AIC:                     770.3
Df Residuals:          461     BIC:                     778.5
Df Model:               1
Covariance Type:        nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	4.1398	0.130	31.865	0.000	3.884	4.395
age	-0.0029	0.003	-1.111	0.267	-0.008	0.002
Omnibus:		14.980	Durbin-Watson:			1.344
Prob(Omnibus):		0.001	Jarque-Bera (JB):			15.952
Skew:		-0.449	Prob(JB):			0.000344
Kurtosis:		2.857	Cond. No.			249.

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Calculate the percentage of visible minorities are tenure professors. Will you say that tenure status differed if teacher was a visible minority?

```
[29]: # Calculate the percentage of visible minorities who are tenure professors
visible_minority_tenure = data[(data['minority'] == 'yes') & (data['tenure'] == 'yes')].shape[0]
visible_minority_total = data[data['minority'] == 'yes'].shape[0]
percentage_visible_minority_tenure = (visible_minority_tenure / visible_minority_total) * 100

print("Percentage of visible minorities who are tenure professors:", percentage_visible_minority_tenure)

# Calculate the percentage of non-visible minorities who are tenure professors
non_visible_minority_tenure = data[(data['minority'] == 'no') & (data['tenure'] == 'yes')].shape[0]
non_visible_minority_total = data[data['minority'] == 'no'].shape[0]
percentage_non_visible_minority_tenure = (non_visible_minority_tenure / non_visible_minority_total) * 100

print("Percentage of non-visible minorities who are tenure professors:", percentage_non_visible_minority_tenure)

# Compare the percentages
if percentage_visible_minority_tenure != percentage_non_visible_minority_tenure:
    print("Tenure status appears to differ based on whether the teacher is a visible minority.")
else:
    print("Tenure status does not appear to differ based on whether the teacher is a visible minority.")
```

Percentage of visible minorities who are tenure professors: 84.375

Percentage of non-visible minorities who are tenure professors:
76.94235588972431

Tenure status appears to differ based on whether the teacher is a visible minority.

Does average age differ by tenure? Produce the means and standard deviations for both tenured and untenured professors.

```
[30]: # Calculate the average and standard deviation of age for tenured professors
tenured_age_mean = data[data['tenure'] == 'yes']['age'].mean()
tenured_age_std = data[data['tenure'] == 'yes']['age'].std()

# Calculate the average and standard deviation of age for untenured professors
untenured_age_mean = data[data['tenure'] == 'no']['age'].mean()
untenured_age_std = data[data['tenure'] == 'no']['age'].std()

# Print the results
print("Average age of tenured professors:", tenured_age_mean)
print("Standard deviation of age for tenured professors:", tenured_age_std)

print("\nAverage age of untenured professors:", untenured_age_mean)
print("Standard deviation of age for untenured professors:", untenured_age_std)
```

Average age of tenured professors: 47.850415512465375
Standard deviation of age for tenured professors: 10.420055773692855

Average age of untenured professors: 50.18627450980392
Standard deviation of age for untenured professors: 6.946372216733221

2 CODE BY ADITYA KUMAR (UI22CS03)